

Redrob Candidate Ranker



Problem

Recruiters need a trusted shortlist for a Senior AI Engineer founding-team role.

The JD explicitly warns against simple keyword matching, so the system ranks for real production retrieval/ranking experience, shipping judgment, platform activity, and logistics.



Target Candidate

- 5-9 years preferred, with flexibility for strong signals.
 - Applied ML/AI engineer with production embeddings, vector search, hybrid retrieval, ranking, and evaluation experience.
 - Product-company or marketplace exposure preferred over services-only career history.
 - India location or relocation fit for Pune/Noida hybrid work.
 - Available and responsive on the Redrob platform.
-

Scoring Model

The ranker combines five interpretable components:

- Skills: Python, retrieval, vector databases, ranking/recommendations, evaluation, LLM/fine-tuning, and production systems.
 - Career: current title, shipped search/ranking/recommendation systems, product-company exposure, experience band, and negative fit signals.
 - Behavior: open-to-work, response rate, recent activity, notice period, interview completion, GitHub activity, verification, recruiter saves, and profile views.
 - Logistics: India, Pune/Noida/Delhi NCR/Mumbai/Hyderabad/Bangalore, relocation, and hybrid/onsite/flexible preference.
 - Education: relevant CS/data/AI fields and institution tier.
-

Honeypot Resistance


The model downweights profiles that look strong only by keyword count:

- Expert skills with near-zero duration.
- Very broad skill lists for low experience.
- AI-keyword-heavy non-technical current titles.
- Timeline inconsistencies between total experience and career history.
- Services-only profiles without product or production ML evidence.



Reproducibility

``rank.py`` streams ``candidates.jsonl``, scores candidates deterministically, sorts by score and candidate ID, and writes exactly 100 rows with grounded 1-2 sentence reasoning. It uses only Python standard library modules and makes no network calls.



Submission Artifacts

- ``submission.csv``: required top-100 ranking.
- ``rank.py``: complete ranking code.
- ``submission_metadata.yaml``: portal metadata mirror.
- ``README.md``: reproduce and validation commands.
- ``approach_deck.pdf``: this methodology as a PDF.

